

docs_chain — System Architecture

Revision: 1 **Last modified:** 2026-05-29T00:00:00Z **Status:** Design documentation — see per-section status tags (IMPLEMENTED vs PLANNED (Phase N)) **Authority:** Operator mandate 2026-05-29 (docs_chain initiative) **Design provenance:** authoritative Phase-0 DESIGN / RESEARCH / PLAN live in the consuming project research tree (docs/research/docs_chain/); this document is the self-contained specification.

Status legend

Per the anti-bluff no-guessing mandate (§11.4.6), every architectural component below carries an explicit implementation-status tag:

- **IMPLEMENTED** — code exists in docs_chain/*.go and behaves as described.
- **PLANNED (Phase N)** — designed in the Phase 0 artefacts, scheduled for Phase N of PLAN.md, NOT yet built.

At the time of writing, the Go package tree under docs_chain/ contains only go.mod and an internal/ scaffold (Phase 1 owns the engine). Therefore **every behavioural component in this document is PLANNED** unless a future revision of this file marks it IMPLEMENTED. This document describes the DESIGNED system; it does not claim working behaviour that is not yet built.

1. What docs_chain is

docs_chain is a universal, Go-implemented, **bidirectional document-and-database dependency-propagation engine**. When any member of a registered chain changes — a Markdown source, an HTML/PDF export, or a SQLite database — docs_chain detects the change by **content hash** and propagates it through

every connected member in every declared direction, regenerating and re-exporting as needed, so that no tracked artefact can fall out of sync.

The engine is the mechanical successor to the project's ad-hoc sync scripts (`sync_issues_docs.sh`, `sync_all_markdown_exports.sh`, `generate_issues_summary.sh`, `generate_fixed_summary.sh`, `sync_asset_player_status.sh`, and `siblings`). It supersedes them without breaking them during migration: each script is wrapped as an `exec: transform` until its logic is migrated into a builtin (see §10).

The recommended conceptual model (from `RESEARCH.md` §9) is **Salsa-style content-hashed incremental recomputation over a DAG**, with Kahn topological ordering, early-cutoff, declared-authority bidirectional edges, and atomic-rename + SQLite-transaction commit.

2. The DAG model

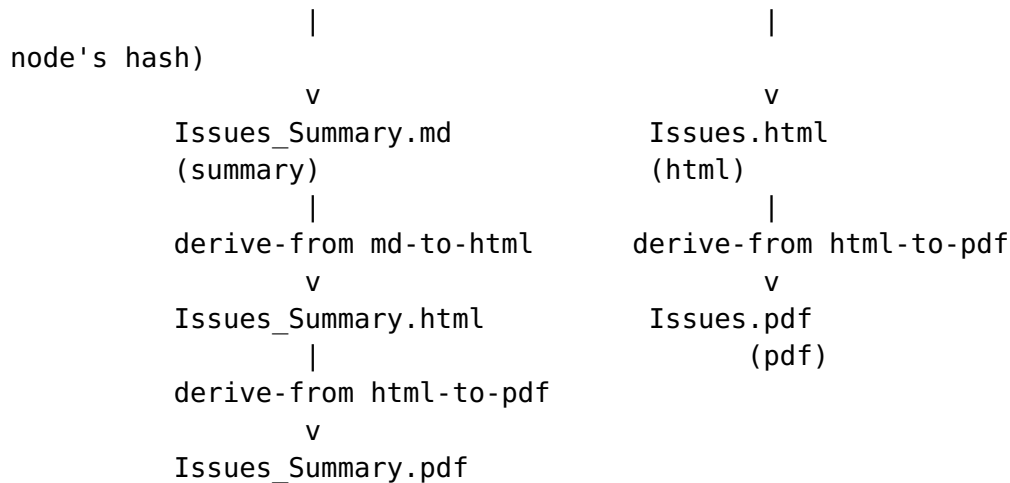
Status: PLANNED (Phase 1 — graph; Phase 2 — node adapters).

A **context** owns one or more **chains**. A chain is a directed graph of **nodes** connected by **edges**.

2.1 Node kinds

Each node kind is a typed adapter (DESIGN §6). A node stores its last-known content hash in the chain state file.

Node kind	Role	Notes
markdown	canonical .md source	input node (Salsa "square")
html	pandoc-derived export	derived
pdf	weasyprint-derived export	derived
sqlite	a database	input AND derived (bidirectional)
summary	generated Markdown digest	e.g. <code>Issues_Summary.md</code>
status	§11.4.45 status doc	derived
status_summary	§11.4.56 two-audience digest	derived
fingerprint		



3. Content-hash change detection (NOT mtime)

Status: PLANNED (Phase 1 — internal/hash).

Change detection keys on **content hash**, never on modification time. This directly encodes the §11.4.86 forensic lesson: git checkout resets mtime, so an mtime-based fingerprint silently misses real content drift and falsely flags unchanged files. Build-tooling practice (Bazel / Buck2 content-addressing, RESEARCH.md §3) corroborates the same choice.

- Per node: sha256(canonicalised content).
- Per multi-member fingerprint node (roster/corpus): sha256 of the **sorted member list** — the existing §11.4.86 pattern.
- mtime is used ONLY as a cheap pre-filter to avoid re-reading unchanged files; it is never the authority.

A node is **dirty** when its freshly-computed hash differs from the hash recorded in state.json.

4. Early-cutoff incremental recompute

Status: PLANNED (Phase 1 — early-cutoff prune; Phase 3 — orchestrator).

Early-cutoff is the load-bearing Salsa feature (RESEARCH.md §1). At every derived node:

1. Recompute the node from its sources via its transform.
2. Hash the result.

3. If the new hash equals the stored hash, **prune the subtree** — no downstream work, no write.

This means `md → db → summary → html → pdf` does NOT cascade work when an edit is a no-op at a downstream layer (for example, a Markdown whitespace-only edit that the summary generator normalises away stops the chain at the summary node). Early-cutoff is also what prevents the bidirectional `md ↔ db` edge from oscillating (§5).

5. Bidirectional sync-edge authority + conflict semantics

Status: PLANNED (Phase 3 — sync-edge authority resolution).

A naïve `md ↔ db` edge is a 2-cycle that would loop forever. `docs_chain` resolves it with a declared-authority algorithm (DESIGN §3):

1. On a change event, hash ALL nodes and compare to stored hashes → the **dirty set**.
2. For a sync pair, if **exactly one** side is dirty, that side is the **source of truth for this run**; the other side is regenerated from it via the appropriate transform (`md→db` or `db→md`).
3. If **both** sides are dirty (a concurrent edit on each side), `docs_chain` **stops and emits a conflict** — exit non-zero, no writes. The operator resolves it (§11.4.66). No silent merge, no guessing (§11.4.6).
4. After regenerating the non-authoritative side, its new hash is recorded. Because the hash now matches what the transform produced, **early-cutoff** prevents it from being re-flagged dirty on the next pass → no loop.

The `authority:` field in the edge config (§7) names the default source of truth for a sync pair, used to disambiguate the "exactly one dirty" case where the dirty side equals the declared authority, and used in diagnostics.

5.1 Bidirectional `md↔db` sync (Mermaid sequence)

```
sequenceDiagram
    participant U as Operator / editor
    participant DC as docs_chain sync
    participant H as hash(all nodes)
    participant R as sync resolver
    participant T as transform
    participant S as state.json
```

```

U->>DC: edit Issues.md (md dirty)
DC->>H: compute hashes
H->>DC: dirty = {Issues.md}
DC->>R: resolve sync pair (md, db)
Note over R: exactly one dirty (md)<br/>md is source of truth
R->>T: md→db transform
T->>R: workable_items.db regenerated
R->>H: re-hash db
Note over R: db hash now matches transform output<br/>early-
cutoff: db not re-flagged
DC->>T: derive-from md→summary→html→pdf
DC->>S: record new hashes (atomic commit)

```

5.2 Both-dirty conflict (ASCII)

```

edit A on Issues.md          edit B on workable_items.db
      \                      /
      v                      v
dirty = { Issues.md , workable_items.db }
      |
      | sync resolver: BOTH dirty
      |
      v
CONFLICT -> exit 2, ZERO writes
(operator resolves per §11.4.66;
no silent merge, no guess per §11.4.6)

```

6. Kahn topological propagation

Status: PLANNED (Phase 1 — Kahn topo-sort + cycle detection).

Propagation runs in deterministic topological order using Kahn's algorithm (RESEARCH.md §5):

- The **derive-from sub-graph MUST be acyclic**. docs_chain runs Kahn's algorithm at load; any residual nodes (a cycle) ⇒ config error and the engine refuses to run (exit 4).
- **sync pairs are collapsed to their authority** before the topo-sort, so the effective propagation graph is always a DAG.
- After the authority of each sync pair is resolved (§5), derive-from edges run from that authority outward (md → summary → html → pdf).

6.1 Loop / cycle prevention (defence in depth)

1. **Load-time Kahn check** — a cycle in derive-from ⇒ refuse to run.

2. **sync collapse** — bidirectional edges never enter the topo-sort as cycles.
 3. **Early-cutoff** (§4) — prevents md↔db oscillation.
 4. **Per-run visited-set + max-iteration guard** — belt-and-suspenders against any unexpected oscillation: fail loudly, never spin.
-

7. Propagation-flow overview (Mermaid)

flowchart TD

```
START([docs_chain sync context]) --> LOAD[Load context  
YAML<br/>build graph]  
LOAD --> KAHN{Kahn cycle<br/>check on<br/>derive-from?}  
KAHN -->|cycle| ERR4([exit 4: config error])  
KAHN -->|acyclic| HASH[Hash ALL nodes<br/>vs state.json]  
HASH --> DIRTY{Any dirty?}  
DIRTY -->|none| EXIT0OK([exit 0: in-sync])  
DIRTY -->|yes| SYNC{sync pair<br/>both dirty?}  
SYNC -->|both| ERR2([exit 2: conflict<br/>no writes])  
SYNC -->|one or none| AUTH[Resolve sync authority<br/>  
>regenerate other side]  
AUTH --> TOPO[Kahn topo order<br/>from authority outward]  
TOPO --> RECOMP[Recompute derived node]  
RECOMP --> CUT{Hash changed<br/>vs stored?}  
CUT -->|no| PRUNE[Early-cutoff:<br/>prune subtree]  
CUT -->|yes| STAGE[Stage temp artefact]  
PRUNE --> MORE{More nodes?}  
STAGE --> MORE  
MORE -->|yes| RECOMP  
MORE -->|no| COMMIT[Atomic commit:<br/>rename temps + DB txn +  
state.json]  
COMMIT --> EXIT0([exit 0: applied])  
RECOMP -->|transform error| ERR3([exit 3: rollback,<br/>no  
live changes])
```

8. Atomicity & crash-safety (composes with §9.2)

Status: PLANNED (Phase 3 — staging + atomic-rename commit + rollback).

A propagation run is all-or-nothing (DESIGN §5, RESEARCH.md §6):

- Each derived artefact is written to <file>.docs_chain.tmp in the same directory, then fsync'd.

- SQLite mutations run inside **one** transaction; `wal_checkpoint(TRUNCATE)` runs before commit (§11.4.95), so the transient `.db-wal` / `.db-shm` sidecars are safely discardable.
- **Commit phase:** atomic `rename(2)` of every staged temp over its live artefact, then commit the SQLite transaction, then write `state.json`.
- **Rollback:** any error before commit $\Rightarrow$ delete all temps, abort the transaction, and leave every live artefact and the DB byte-identical to the pre-run state. Exit 3.
- For destructive contexts an optional pre-run hardlinked `.git` backup hook composes with the §9.2 zero-risk data-safety protocol.

Crash-safety claim verified by the Phase 3 / Phase 5 chaos tests (§11.4.85): SIGKILL mid-commit leaves live artefacts byte-identical to the pre-run snapshot (`recovery_trace.log` is the captured evidence).

9. SQLite-node integration (§11.4.93)

Status: PLANNED (Phase 2 — `sqlite $\leftrightarrow$ md` pluggable transform; Phase 3 — `txn/WAL` commit).

The `sqlite` node is the bidirectional pivot of the §11.4.93 single source of truth (`docs/workable_items.db`, tracked in git per §11.4.95). Two pluggable transforms bind it to its Markdown view:

- `md $\rightarrow$ sqlite` — parse Markdown trackers into DB rows.
- `sqlite $\rightarrow$ md` — regenerate the Markdown trackers from DB rows.

During migration both transforms shell out to the existing §11.4.93 `workable-items` Go binary (DESIGN §6) so there is zero rewrite risk. The DB participates in the same atomic-commit phase as the file artefacts (§8): the transaction is committed only after every staged file rename succeeds, so a partial run never leaves the DB ahead of (or behind) the Markdown.

9.1 Multi-context overview (Mermaid)

A consuming project registers several independent contexts; each is a self-contained chain and `docs_chain sync <context>` touches only that context's nodes.

```
graph TB
    subgraph issues[context: issues]
        I_DB[(workable_items.db)] <--> I_MD[Issues.md]
        I_MD --> I_SUM[Issues_Summary.md] --> I_EXP[html + pdf]
        I_MD --> I_EXP2[html + pdf]
    end
```



```

    subgraph fixed[context: fixed]
        F_MD[Fixed.md] --> F_SUM[Fixed_Summary.md] --> F_EXP[html
+ pdf]
        F_MD --> F_EXP2[html + pdf]
    end
    subgraph status[context: status]
        S_MD[Status.md] --> S_SUM[Status_Summary.md] -->
S_EXP[html + pdf]
    end
    subgraph asset[context: asset_player]
        A_FP[roster fingerprint] --> A_ST[players/Status.md] -->
A_EXP[html + pdf]
    end
    CLI{{docs_chain sync --all}} --> issues
    CLI --> fixed
    CLI --> status
    CLI --> asset

```

9.2 Multi-context overview (ASCII)

```

docs_chain sync --all
|
+-- context: issues          -> workable_items.db <-> Issues.md
-> summary -> html/pdf
+-- context: fixed          -> Fixed.md -> Fixed_Summary.md ->
html/pdf
+-- context: status        -> Status.md -> Status_Summary.md
-> html/pdf
+-- context: asset_player  -> roster fingerprint -> players/
Status.md -> html/pdf
|
each context is independent; `sync issues` touches ONLY the
issues chain.

```

10. Supersession without breakage (migration model)

Status: PLANNED (Phase 7 — ATMOSphere wiring + retire ad-hoc scripts).

Each ad-hoc sync script is wrapped as an exec: transform (DESIGN §8). Phase 7 registers the ATMOSphere contexts pointing at the EXISTING scripts: behaviour is identical, but now orchestrated, content-hashed, atomic, and conflict-aware. Once a context is green for N cycles, the script's internal re-export / parity logic migrates into docs_chain builtins and the script retires to a thin shim that calls docs_chain sync <context>.

See [USE_CASE_CATALOGUE.md](#) for the full script-to-context supersession table.

11. Go package layout

Status: PLANNED (Phases 1-4).

```
docs_chain/
  cmd/docs_chain/main.go      # CLI: sync | watch | verify | graph
  | doctor
  internal/graph/             # DAG, Kahn topo-sort, cycle detect,
early-cutoff
  internal/hash/              # sha256 canonical content + member-
list fingerprint
  internal/node/              # node kinds + adapters (uses vasic-
digital/Document)
  internal/transform/         # builtin + exec transforms
  internal/orchestrator/      # dirty-set, sync-resolution,
staging, atomic commit, rollback
  internal/config/            # YAML context loader + validation
  internal/watch/             # fsnotify daemon (uses vasic-
digital/Watcher)
  internal/state/             # state.json read/write (§11.4.77
regen)
```

The engine reuses three catalogued vasic-digital modules (RESEARCH.md §8): Document (node-content abstraction), Watcher (watch daemon), and Lazy (memoised node evaluation). No existing dependency-graph / incremental-propagation engine was found, so the engine itself is original work extending those three modules (§11.4.8).

12. CLI / daemon interface

Status: PLANNED (Phase 4 — CLI/daemon).

Command	Behaviour	Exit codes
docs_chain sync <context>	one-shot: detect dirty, resolve, propagate atomically, update state	0 in-sync/applied · 2 conflict · 3 transform-fail · 4 cycle/config-error
docs_chain sync --all	every registered context	as above
docs_chain verify <context>	read-only: report drift, write nothing (the pre-	0 in-sync · non-zero drift

Command	Behaviour	Exit codes
	build-gate hook; the --check-only successor)	
docs_chain watch [--context X]	fsnotify daemon, debounced, runs sync on settle	long-running
docs_chain graph <context>	emit the DAG (DOT/JSON) for audit	0
docs_chain doctor	validate all contexts + state integrity	0 healthy

13. Anti-bluff binding

Status: PLANNED (Phase 5 — comprehensive test suite).

Every transform run captures evidence (input hashes, output hashes, rename log) into qa-results/docs_chain/<run-id>/ (§11.4.69 derived- artefact feature class). verify is the deterministic (§11.4.50) sink-side check that a consuming project's pre-build gate invokes. The full four-layer test coverage per component (pre-build gate, packaging check, runtime/e2e test, paired §1.1 meta-test mutation) is specified in PLAN.md Phase 5.