

ARCHITECTURE

Docs Chain — System Architecture

Revision: 3 **Last modified:** 2026-05-29T12:00:00Z **Status:** Design documentation — see per-section status tags (IMPLEMENTED vs PLANNED (Phase N)). Phases 1–3 are now IMPLEMENTED + tested (internal/{hash,graph,adapter,orchestrator}); Phases 4–7 remain PLANNED.

Authority: Operator mandate 2026-05-29 (Docs Chain initiative) **Design provenance:** authoritative Phase-0 DESIGN / RESEARCH / PLAN live in the consuming project research tree (docs/research/docs_chain/); this document is the self-contained specification.

Status legend

Per the anti-bluff no-guessing mandate (§11.4.6), every architectural component below carries an explicit implementation-status tag:

- **IMPLEMENTED** — code exists in docs_chain/*.go and behaves as described.
- **PLANNED (Phase N)** — designed in the Phase 0 artefacts, scheduled for Phase N of PLAN.md, NOT yet built.

As of Revision 2, the Go package tree under internal/ implements Phases 1–3: internal/hash + internal/graph (engine), internal/adapter (node adapters + transforms), and internal/orchestrator (atomic propagation). go test ./... passes. Components owned by Phases 4–7 (the cmd/ CLI/daemon, per-context YAML config, the beyond-package test suite, constitution-submodule integration) remain **PLANNED** and are tagged as such per section. This document describes the DESIGNED system and does not claim working behaviour that is not yet built (§11.4.6).

1. What Docs Chain is

Docs Chain is a universal, Go-implemented, **bidirectional document-and-database dependency-propagation engine**. When any member of a registered chain changes — a Markdown source, an HTML/PDF export, or a SQLite database — Docs Chain

detects the change by **content hash** and propagates it through every connected member in every declared direction, regenerating and re-exporting as needed, so that no tracked artefact can fall out of sync.

The engine is the mechanical successor to the project’s ad-hoc sync scripts (`sync_issues_docs.sh`, `sync_all_markdown_exports.sh`, `generate_issues_summary.sh`, `generate_fixed_summary.sh`, `sync_asset_player_status.sh`, and `siblings`). It supersedes them without breaking them during migration: each script is wrapped as an `exec: transform` until its logic is migrated into a builtin (see §10).

The recommended conceptual model (from `RESEARCH.md` §9) is **Salsa-style content-hashed incremental recomputation over a DAG**, with Kahn topological ordering, early-cutoff, declared-authority bidirectional edges, and atomic-rename + SQLite-transaction commit.

2. The DAG model

Status: IMPLEMENTED (Phase 1 — graph in `internal/graph`; Phase 2 — node adapters in `internal/adapters`).

A **context** owns one or more **chains**. A chain is a directed graph of **nodes** connected by **edges**.

2.1 Node kinds

Each node kind is a typed adapter (`DESIGN` §6). A node stores its last-known content hash in the chain state file.

Node kind	Role	Notes
markdown	canonical <code>.md</code> source	input node (Salsa “square”)
html	pandoc-derived export	derived
pdf	weasyprint-derived export	derived
sqlite	a database	input AND derived (bidirectional)
summary	generated Markdown digest	e.g. <code>Issues_Summary.md</code>
status	§11.4.45 status doc	derived
status_summary	§11.4.56 two-audience digest	derived
fingerprint	roster/corpus sidecar	§11.4.86 member-list hash

2.2 Edge kinds

- **derive-from** (one-way): the target is regenerated from one or more sources via a transform. Examples: `html ← markdown`, `pdf ← html`, `summary ← markdown`.
- **sync** (bidirectional): two nodes are mutually authoritative views of the same data. Example: `markdown ↔ sqlite` (§11.4.93). A sync edge is NOT a literal 2-cycle; it is a **paired relation with a declared authority plus two transforms** (`md→db` and `db→md`). See §5.


2.3 DAG example (Mermaid)


The canonical issues chain (§11.4.93 / §11.4.12):

```
graph LR
    DB[(workable_items.db<br/>sqlite)]
    MD[Issues.md<br/>markdown]
    SUM[Issues_Summary.md<br/>summary]
    HTML[Issues.html<br/>html]
    PDF[Issues.pdf<br/>pdf]

    DB <-->|sync<br/>authority: db| MD
    MD -->|derive-from<br/>gen-issues-summary| SUM
    MD -->|derive-from<br/>md-to-html| HTML
    HTML -->|derive-from<br/>html-to-pdf| PDF
    SUM -->|derive-from<br/>md-to-html| HTML
    SUMHTML[Issues_Summary.html<br/>html]
    SUMHTML -->|derive-from<br/>html-to-pdf| PDF
    SUMPDF[Issues_Summary.pdf<br/>pdf]
```

2.4 DAG example (ASCII)





3. Content-hash change detection (NOT mtime)

Status: PLANNED (Phase 1 — internal/hash).

Change detection keys on **content hash**, never on modification time. This directly encodes the §11.4.86 forensic lesson: `git checkout` resets mtime, so an mtime-based fingerprint silently misses real content drift and falsely flags unchanged files. Build-tooling practice (Bazel / Buck2 content-addressing, RESEARCH.md §3) corroborates the same choice.

- Per node: `sha256(canonicalised content)`.
- Per multi-member fingerprint node (roster/corpus): `sha256` of the **sorted member list** — the existing §11.4.86 pattern.
- mtime is used **ONLY** as a cheap pre-filter to avoid re-reading unchanged files; it is never the authority.

A node is **dirty** when its freshly-computed hash differs from the hash recorded in `state.json`.

4. Early-cutoff incremental recompute

Status: PLANNED (Phase 1 — early-cutoff prune; Phase 3 — orchestrator).

Early-cutoff is the load-bearing Salsa feature (RESEARCH.md §1). At every derived node:

1. Recompute the node from its sources via its transform.
2. Hash the result.
3. If the new hash equals the stored hash, **prune the subtree** — no downstream work, no write.

This means `md → db → summary → html → pdf` does NOT cascade work when an edit is a no-op at a downstream layer (for example, a Markdown whitespace-only edit that the summary generator normalises away stops the chain at the summary node). Early-cutoff is also what prevents the bidirectional `md` `db` edge from oscillating (§5).

5. Bidirectional sync-edge authority + conflict semantics

Status: IMPLEMENTED (Phase 1 — `graph.ResolveSync` authority/conflict detection; Phase 3 — `orchestrator.Run` surfaces `ConflictError` with zero writes).

A naïve md $\leftrightarrow$ db edge is a 2-cycle that would loop forever. Docs Chain resolves it with a declared-authority algorithm (DESIGN §3):

1. On a change event, hash ALL nodes and compare to stored hashes → the **dirty set**.
2. For a sync pair, if **exactly one** side is dirty, that side is the **source of truth for this run**; the other side is regenerated from it via the appropriate transform (md→db or db→md).
3. If **both** sides are dirty (a concurrent edit on each side), Docs Chain **stops and emits a conflict** — exit non-zero, no writes. The operator resolves it (§11.4.66). No silent merge, no guessing (§11.4.6).
4. After regenerating the non-authoritative side, its new hash is recorded. Because the hash now matches what the transform produced, **early-cutoff** prevents it from being re-flagged dirty on the next pass → no loop.

The `authority:` field in the edge config (§7) names the default source of truth for a sync pair, used to disambiguate the “exactly one dirty” case where the dirty side equals the declared authority, and used in diagnostics.

5.1 Bidirectional md $\leftrightarrow$ db sync (Mermaid sequence)

sequenceDiagram

```
participant U as Operator / editor
participant DC as docs_chain sync
participant H as hash(all nodes)
participant R as sync resolver
participant T as transform
participant S as state.json
```

```
U->>DC: edit Issues.md (md dirty)
```

```
DC->>H: compute hashes
```

```
H-->>DC: dirty = {Issues.md}
```

```
DC->>R: resolve sync pair (md, db)
```

```
Note over R: exactly one dirty (md)<br/>md is source of truth
```

```
R->>T: md→db transform
```

```
T-->>R: workable_items.db regenerated
```

```
R->>H: re-hash db
```

```
Note over R: db hash now matches transform output<br/>early-cutoff: db not re-flagged
```

```
DC->>T: derive-from md→summary→html→pdf
```

```
DC->>S: record new hashes (atomic commit)
```

5.2 Both-dirty conflict (ASCII)

```
edit A on Issues.md          edit B on workable_items.db
      \                      /
      v                      v
dirty = { Issues.md , workable_items.db }
      |
      sync resolver: BOTH dirty
      |
      v
CONFLICT -> exit 2, ZERO writes
(operator resolves per §11.4.66;
no silent merge, no guess per §11.4.6)
```

6. Kahn topological propagation

Status: IMPLEMENTED (Phase 1 — Kahn topo-sort + cycle detection in `graph.TopoOrder` / `graph.Validate`; Phase 3 — `orchestrator.Run cycle-guard`).

Propagation runs in deterministic topological order using Kahn's algorithm (`RESEARCH.md` §5):

- The **derive-from sub-graph MUST be acyclic**. Docs Chain runs Kahn's algorithm at load; any residual nodes (a cycle) $\Rightarrow$ config error and the engine refuses to run (exit 4).
- **sync pairs are collapsed to their authority** before the topo-sort, so the effective propagation graph is always a DAG.
- After the authority of each sync pair is resolved (§5), derive-from edges run from that authority outward (md $\rightarrow$ summary $\rightarrow$ html $\rightarrow$ pdf).

6.1 Loop / cycle prevention (defence in depth)

1. **Load-time Kahn check** — a cycle in derive-from $\Rightarrow$ refuse to run.
 2. **sync collapse** — bidirectional edges never enter the topo-sort as cycles.
 3. **Early-cutoff** (§4) — prevents md $\leftrightarrow$ db oscillation.
 4. **Per-run visited-set + max-iteration guard** — belt-and-suspenders against any unexpected oscillation: fail loudly, never spin.
-

7. Propagation-flow overview (Mermaid)

flowchart TD

```
START([docs_chain sync context]) --> LOAD[Load context  
YAML<br/>build graph]  
LOAD --> KAHN{Kahn cycle<br/>check on<br/>derive-from?}  
KAHN -->|cycle| ERR4([exit 4: config error])  
KAHN -->|acyclic| HASH[Hash ALL nodes<br/>vs state.json]  
HASH --> DIRTY{Any dirty?}  
DIRTY -->|none| EXIT00K([exit 0: in-sync])
```

```

DIRTY -->|yes| SYNC{sync pair<br/>both dirty?}
SYNC -->|both| ERR2([exit 2: conflict<br/>no writes])
SYNC -->|one or none| AUTH[Resolve sync authority<br/>
>regenerate other side]
AUTH --> TOP0[Kahn topo order<br/>from authority outward]
TOP0 --> RECOMP[Recompute derived node]
RECOMP --> CUT{Hash changed<br/>vs stored?}
CUT -->|no| PRUNE[Early-cutoff:<br/>prune subtree]
CUT -->|yes| STAGE[Stage temp artefact]
PRUNE --> MORE{More nodes?}
STAGE --> MORE
MORE -->|yes| RECOMP
MORE -->|no| COMMIT[Atomic commit:<br/>rename temps + DB txn +
state.json]
COMMIT --> EXIT0([exit 0: applied])
RECOMP -->|transform error| ERR3([exit 3: rollback,<br/>no
live changes])

```

8. Atomicity & crash-safety (composes with §9.2)

Status: IMPLEMENTED (Phase 3 — orchestrator.Run in-memory staging + withhold-until-success + per-file atomic temp-then-rename + rollback on any transform error). The optional pre-run hardlinked .git backup hook and the SIGKILL-mid-commit chaos verification remain PLANNED (Phase 5).

A propagation run is all-or-nothing (DESIGN §5, RESEARCH.md §6):

- Each derived artefact is written to <file>.docs_chain.tmp in the same directory, then fsync'd.
- SQLite mutations run inside **one** transaction; wal_checkpoint(TRUNCATE) runs before commit (§11.4.95), so the transient .db-wal / .db-shm sidecars are safely discardable.
- **Commit phase:** atomic rename(2) of every staged temp over its live artefact, then commit the SQLite transaction, then write state.json.
- **Rollback:** any error before commit ⇒ delete all temps, abort the transaction, and leave every live artefact and the DB byte-identical to the pre-run state. Exit 3.
- For destructive contexts an optional pre-run hardlinked .git backup hook composes with the §9.2 zero-risk data-safety protocol.

Crash-safety claim verified by the Phase 3 / Phase 5 chaos tests (§11.4.85): SIGKILL mid-commit leaves live artefacts byte-identical to the pre-run snapshot (recovery_trace.log is the captured evidence).

As-built (Revision 2, Phase 3 — IMPLEMENTED). orchestrator.Run implements the staging + withhold + rollback contract: every regenerated output is buffered in-memory and live files are touched only after the recompute fully succeeds; any transform error returns Status: rolled-back with zero files written (verified by TestRun_TransformError_RollsBack). Each FileAdapter.Write is itself atomic (temp-file + Sync + rename). The SQLiteAdapter.Write runs its mutation inside a single database/sql transaction (committed or rolled back as a unit). What remains PLANNED (Phase 5): the unified cross-artefact commit phase

that interleaves the SQLite transaction Commit *between* the file renames and a state.json write (the current build commits each adapter independently in deterministic order and reports a partial Committed list honestly if an OS-level write fails mid-drain), the wal_checkpoint(TRUNCATE) step, the state.json baseline persistence (Phase 3 keeps the baseline in-memory via graph.CommitHashes), and the SIGKILL-mid-commit chaos verification.

9. SQLite-node integration (§11.4.93)

Status: IMPLEMENTED (Phase 2 — adapter.SQLiteAdapter: canonical row-dump content for hashing + transaction-wrapped Apply for the md→db write side, pure-Go modernc.org/sqlite, no cgo). The unified cross-artefact WAL-checkpoint commit phase remains PLANNED (Phase 5, see §8 as-built).

The sqlite node is the bidirectional pivot of the §11.4.93 single source of truth (docs/workable_items.db, tracked in git per §11.4.95). Two pluggable transforms bind it to its Markdown view:

- md→sqlite — parse Markdown trackers into DB rows.
- sqlite→md — regenerate the Markdown trackers from DB rows.

During migration both transforms shell out to the existing §11.4.93 workable-items Go binary (DESIGN §6) so there is zero rewrite risk.

As-built (Revision 2). The SQLiteAdapter hashes a **canonical dump** of a caller-supplied deterministic ORDER BY query — NOT the raw .db page bytes — so logically-identical databases collide regardless of insertion order or WAL/page noise (verified by TestSQLiteAdapter_CanonicalDumpDeterminism). The write side runs the caller's Apply mutation inside one database/sql transaction (TestSQLiteAdapter_WriteApplyTransaction); a node with no Apply func is read-only and rejects Write. The DumpQuery (the documented query hook) is mandatory. The cross-artefact commit ordering between the DB transaction and the file renames is owned by the Phase-5 unified commit phase (§8 as-built); the current build commits each adapter independently.

9.1 Multi-context overview (Mermaid)

A consuming project registers several independent contexts; each is a self-contained chain and docs_chain sync <context> touches only that context's nodes.

```
graph TB
    subgraph issues[context: issues]
        I_DB[(workable_items.db)] <--> I_MD[Issues.md]
        I_MD --> I_SUM[Issues_Summary.md] --> I_EXP[html + pdf]
        I_MD --> I_EXP2[html + pdf]
    end
    subgraph fixed[context: fixed]
        F_MD[Fixed.md] --> F_SUM[Fixed_Summary.md] --> F_EXP[html + pdf]
        F_MD --> F_EXP2[html + pdf]
    end
```



```

    subgraph status[context: status]
        S_MD[Status.md] --> S_SUM[Status_Summary.md] -->
S_EXP[html + pdf]
    end
    subgraph asset[context: asset_player]
        A_FP[roster fingerprint] --> A_ST[players/Status.md] -->
A_EXP[html + pdf]
    end
    CLI{{docs_chain sync --all}} --> issues
    CLI --> fixed
    CLI --> status
    CLI --> asset

```

9.2 Multi-context overview (ASCII)

```

docs_chain sync --all
|
+-- context: issues          -> workable_items.db <-> Issues.md
-> summary -> html/pdf
+-- context: fixed          -> Fixed.md -> Fixed_Summary.md ->
html/pdf
+-- context: status         -> Status.md -> Status_Summary.md
-> html/pdf
+-- context: asset_player   -> roster fingerprint -> players/
Status.md -> html/pdf
|
each context is independent; `sync issues` touches ONLY the
issues chain.

```

10. Supersession without breakage (migration model)

Status: PLANNED (Phase 7 — ATMOSphere wiring + retire ad-hoc scripts).

Each ad-hoc sync script is wrapped as an `exec: transform` (DESIGN §8). Phase 7 registers the ATMOSphere contexts pointing at the EXISTING scripts: behaviour is identical, but now orchestrated, content-hashed, atomic, and conflict-aware. Once a context is green for N cycles, the script's internal re-export / parity logic migrates into Docs Chain builtins and the script retires to a thin shim that calls `docs_chain sync <context>`.

See [USE_CASE_CATALOGUE.md](#) for the full script-to-context supersession table.

11. Go package layout

Status: PLANNED (Phases 1–4).

```
docs_chain/
  cmd/docs_chain/main.go      # CLI: sync | watch | verify | graph
  | doctor
  internal/graph/             # DAG, Kahn topo-sort, cycle detect,
early-cutoff
  internal/hash/               # sha256 canonical content + member-
list fingerprint
  internal/node/               # node kinds + adapters (uses vasic-
digital/Document)
  internal/transform/          # builtin + exec transforms
  internal/orchestrator/       # dirty-set, sync-resolution,
staging, atomic commit, rollback
  internal/config/             # YAML context loader + validation
  internal/watch/              # fsnotify daemon (uses vasic-
digital/Watcher)
  internal/state/              # state.json read/write (§11.4.77
regen)
```

The engine reuses three catalogued vasic-digital modules (RESEARCH.md §8): Document (node-content abstraction), Watcher (watch daemon), and Lazy (memoised node evaluation). No existing dependency-graph / incremental-propagation engine was found, so the engine itself is original work extending those three modules (§11.4.8).

12. CLI / daemon interface

Status: PLANNED (Phase 4 — CLI/daemon).

Command	Behaviour	Exit codes
docs_chain sync <context>	one-shot: detect dirty, resolve, propagate atomically, update state	0 in-sync/ applied · 2 conflict · 3 transform-fail · 4 cycle/config- error
docs_chain sync --all	every registered context	as above
docs_chain verify <context>	read-only: report drift, write nothing (the pre-build- gate hook; the --check- only successor)	0 in-sync · non- zero drift
docs_chain watch [--context X]	fsnotify daemon, debounced,	long-running

Command	Behaviour	Exit codes
	runs sync on settle	
docs_chain graph <context>	emit the DAG (DOT/JSON) for audit	0
docs_chain doctor	validate all contexts + state integrity	0 healthy

13. Anti-bluff binding

Status: PLANNED (Phase 5 — comprehensive test suite).

Every transform run captures evidence (input hashes, output hashes, rename log) into qa-results/docs_chain/<run-id>/ (§11.4.69 derived- artefact feature class). verify is the deterministic (§11.4.50) sink-side check that a consuming project's pre-build gate invokes. The full four-layer test coverage per component (pre-build gate, packaging check, runtime/e2e test, paired §1.1 meta-test mutation) is specified in PLAN.md Phase 5.